

Cite, in one sitting.

What this product is for, how an agent actually buys a citation, how we take it to market without burning Shortlist's publishers, and what is already running in this repo.

To: David Henzel From: the Cite build Owner: shortlist.io • separate brand 18 Aug 2026

Purpose

Cite is an **agent-native marketplace for earned link placements**. A coding agent or a marketing agent states an intent — get `example.com/pricing` cited on fintech sites, \$4,000, eight weeks, conservative risk — and Cite constructs and executes the portfolio: which publishers, in what order, at what cadence, with what anchor-text mix.

It is sold as **infrastructure, not as an agency**. The human buyer never has to see the site list to get a result. Shortlist's existing publisher inventory is the supply. Cite is the demand-side interface that agents can call.

The category is shifting. Buyers increasingly need to be *cited* by ChatGPT, Perplexity, and Google AI Overviews, not only ranked by Google. "Citations" is the growing word; "backlinks" is the shrinking one with the bad smell. The name is one word, one spelling, no hyphen — unambiguous when a model types `claude mcp add cite`.

WHO IT IS FOR

- **Primary buyer:** an AI agent acting for an SEO or growth team, with a budget and a target URL.
- **Also:** agencies and in-house SEO teams who want an API instead of a spreadsheet and a Slack thread.
- **Also:** developers wiring link acquisition into an existing growth stack.

THE POSITIONING THAT MAKES THIS NOT A CATALOG

Selection is the wrong interface. A directory an agent browses is a commodity — any incumbent can ship one in a weekend. CrowdReply already occupies “catalog + MCP”. Cite’s primary verb is intent:

```
create_campaign(target_url, topics[], budget, timeframe, risk_tolerance)
```

An allocator does the portfolio construction. Search still exists for inspection and trust, but it is not the workflow. The allocator — and later the lift data behind it — is the product. The API is how you reach it.

How it will work

End to end, one paid placement looks like this. Humans stay in the fulfilment loop on purpose: publishers are never told the buying is programmatic.

01 The agent states intent

Target URL, topics, budget, timeframe, risk tolerance. No site picking required.

02 Inventory is blind until delivery

Pre-delivery, buyers see anonymized handles: Cite Score, topic, traffic band, price, turnaround. Never the domain, never the owner contact, never the seller price. The domain appears as `published_url` when the link is live. That is why a refund guarantee is required — buying blind has to be rational.

03 The allocator builds the portfolio

In v1 this is composite Cite Score per dollar, with five hard constraints in code: per-site quarterly throttle (default 4), exact-match anchor cap, topical relevance floor, no single placement above 25% of budget, and cadence spread across the timeframe. Later it optimizes observed lift per dollar. Lift is measured from day one and never shown to a customer.

04 Credits, then an order

Prepaid Stripe credits, debited at order creation. The agent submits a finished post (title, markdown body, 1–3 links). Structured errors (`TOPIC_MISMATCH` , `WORD_COUNT_LOW` , `INSUFFICIENT_CREDIT` ...) let the agent fix and resubmit without a human round trip.

05 A person still sends the email

Every order produces a Gmail draft in the team inbox, in the voice of a named human. A person edits and sends. That is permanent, not scaffolding. Agent-to-publisher haggling is out — it is the tell that breaks concealment. Throughput is roughly 30–60 orders per operator per day. That, not inventory, is the revenue ceiling.

06 Guarantee, not outcome pricing

Link live and indexed at T+30, or a full refund. Lost between T+30 and T+90 → prorated credit. Rank and AI-citation movement are measured with no promise attached. Outcome pricing was considered and rejected.

THE THREE NUMBERS THAT MUST NEVER BE MIXED

LAYER	WHAT IT IS	CUSTOMER-FACING?
Guarantee	Live and indexed at T+30, or refund. Contractual. Within Cite’s control.	Yes
Measurement	Rank and AI-citation movement at T+0 / T+30 / T+90.	No promise. No per-customer scoreboard.
Allocation	Uses measurement to pick sites. Invisible. Absent from the contract.	Never

Margin sits between `seller_price` and `listed_price` (default markup ×1.6, editable per site). The seller cut is never published in any API response, at any tier.

How we will market it

Be the thing Claude Code reaches for the way it reaches for Vercel. MCP first, REST underneath. A generous free read tier with no card: **the inventory search is the marketing**, because every exploratory query teaches an agent that Cite is where link supply lives.

Where agents discover us

List the MCP server everywhere agents already look: mcp.so, Smithery, Glama, PulseMCP, and awesome-mcp-servers. The install line is the brand: `claude mcp add cite`.

The trial is free placements

465 sites in the database already cost \$0. Free stays free — no service fee. Agent signup is email-only and instant: anonymous (10 results) → registered (50 results, 10 free placements) → funded via Stripe.

Separate brand, quiet ownership

Launch Cite as its own brand. Do not put “a Shortlist company” in the footer. If Cite loudly markets programmatic placement under the Shortlist name, publishers can connect the two in one search. That burns the inventory. It is an operational argument, not a taste argument.

What we will not say

No per-customer lift dashboard. A number you show becomes a promise, and SEO confounders make that unwinable. Aggregate observational benchmarks are acceptable marketing. Agent-assisted negotiation on our side is margin, never marketed, never exposed.

TWO DIFFERENT MOTIONS, ON PURPOSE

- **Cite** is self-serve, usage-priced, developer-marketed infrastructure. Independently sellable.
- **Shortlist** remains relationship-sold and retainer-priced. Cite routes demand into that inventory; it does not replace the agency.

What is already built

Spec v1 was settled with you on 12 Aug 2026. The working product lives on branch `claude/cite-link-placement-spec-pc4ig4`. Hosted endpoint: `https://cite-mcp.d-henzel.workers.dev`.

STANDING

9,453 publisher domains imported from the Shortlist sheet. 8,108 priced (median \$90). Cite Score computed at import. Blind-placement whitelist enforced in code — domains, contacts, seller prices, and markup cannot leak through the public MCP tools.

STANDING

MCP free read tier: `search_sites` , `get_site` , `estimate` , `inventory_stats` . Connect with `claude mcp add --transport http cite https://cite-mcp.d-henzel.workers.dev/mcp` .

STANDING

Operator console at `/admin` : full private inventory, per-site seller price and markup, listed price recomputed live on D1, link-attribute backfill, add-site, analytics (signups, query volume, unmet demand, inventory readiness).

STANDING

Operator MCP at `/admin/mcp` , plus per-person `cka_...` keys from a Connect tab (pre-filled `claude mcp add` command). Shared operator token retained for agents.

STANDING

Free inventory classified: 22 self-serve (Medium, Hashnode, Substack, Quora, Reddit, LinkedIn...), 378 apply-editorial, 80 link-exchange parked and excluded from search.

IN PROGRESS

Shortlist Context Engine SSO for the console. First live sign-in is blocked on a bad OIDC client secret (`invalid_client`). Break-glass token login is on until that is fixed.

NOT IN VO

Stripe credits, `create_campaign` / paid orders, Gmail-draft fulfilment, verification crawler, refunds, and the lift-observation pipeline. Those are the v1 “one real order, end to end” cut in the spec.

Cite is not Shortlist’s repositioning

Your correction on 13 Aug: Cite is a **separate project that feeds into Shortlist**. Shipping Cite does not answer Shortlist’s own positioning question. Keep the two threads distinct.

Quiet ownership means quiet. Cite uses Shortlist inventory and Shortlist operators. Public copy should never make that easy to reverse-engineer.

Open asks

- **Domain and trademark check for “Cite”** — cheap, and it blocks all public naming.
- **Link-attribute backfill** (dofollow / sponsored) across 9,453 sites — mandatory before launch.
- **Confirm Ahrefs / Moz / Majestic licences** actually include API access, not only UI.
- **Locate the correct OIDC client secret** so Shortlist sign-in works and the token fallback can be turned off.
- **Who drives v1** — agent-driven with you, Shortlist developers, or a hire.

Grounded in Cite spec v1 (settled 12 Aug 2026, decisions through 17 Aug) and the working tree on `claude/cite-link-placement-spec-pc4ig4`. This page is a partner memo, not customer marketing.